

facebook

facebook

Python @ Facebook

Angelo Failla
Production Engineer @ Facebook Dublin
October 11th, 2014 -- PyCon Ireland 2014

Agenda

1

How do we use Python at Facebook?

2

Description of a production system based on Python

3

Questions

How do we use Python at Facebook?



- **Centralized repo for back-end systems**
- **A LOT of code.**
- **All Python code hosted there**
- **3rd most used language
(PHP/Hack => C++ => Python)**
- **Still a LOT of code :P**

ONE BUILD TOOL



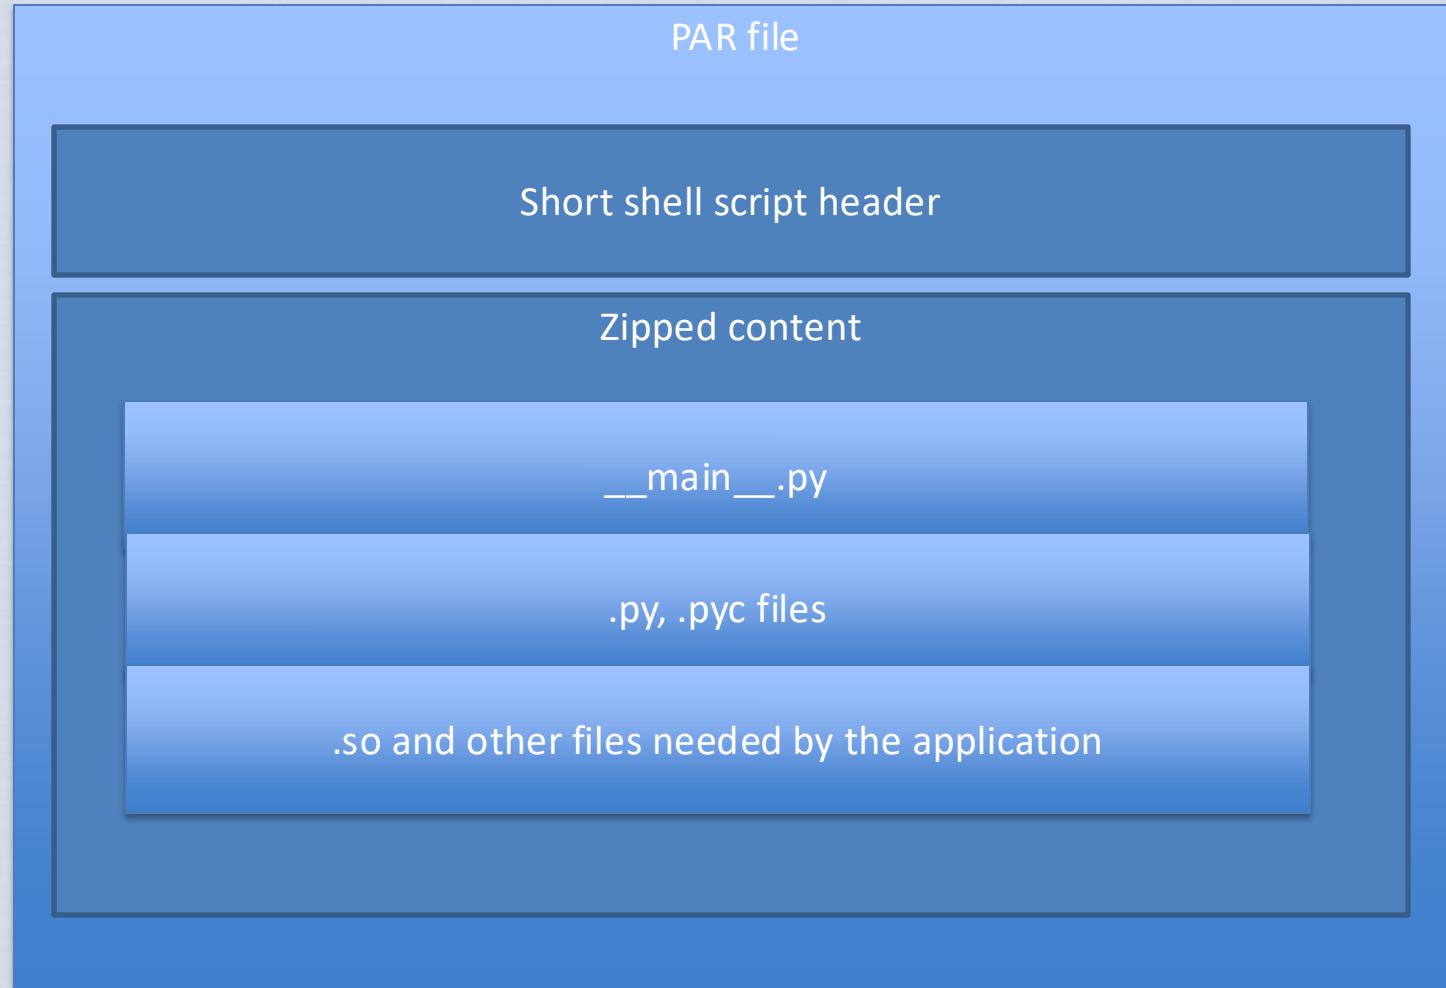
TO RULE THEM ALL

Everything is built from HEAD
C++ binaries are statically linked



<http://cdn.instapop.com/assets/memes/Skeptical%20Baby/4302/thumb.jpeg?1383838705>

How about the same in Python? We use .par files



Thrift (<https://thrift.apache.org/>) in a nutshell

Write `myservice.thrift`

A large blue arrow pointing downwards, indicating the flow from the first step to the second.

Thrift compiler generates language bindings

A large blue arrow pointing downwards, indicating the flow from the second step to the third.

Client/Server app imports generated code

```
# cat
```

```
~/repo/thrift_example/myservice.thrift
```

```
namespace py thrift_example.myservice
```

```
service Calculator {
```

```
    i32 add(1: i32 num1, 2: i32 num2)
```

```
}
```

```
$ thrift --gen py myservice.thrift
```

```
$ tree gen-py/
```

```
gen-py/
```

```
|— __init__.py
```

```
└─ myservice
```

```
    |— Calculator.py
```

```
    |— Calculator-remote
```

```
    |— constants.py
```

```
    |— __init__.py
```

```
    └─ ttypes.py
```

```
# cat ~/fbcode/thrift_example/server.py
# import lines removed, see: http://tinyurl.com/p6a7cuw

1: class CalculatorHandler(Calculator.Iface):
2:     def add(self, num1, num2):
3:         return num1 + num2
4:
5: handler = CalculatorHandler()
6: socket = TSocket.TServerSocket(port=9090)
7: tfactory = TTransport.TBufferedTransportFactory()
8: pfactory = TBinaryProtocol.TBinaryProtocolFactory()
9: server = TServer.TSimpleServer(
10:     handler, socket, tfactory, pfactory)
11: server.serve()
```

```
# cat ~/fbcode/thrift_example/client.py
# import lines removed, see: http://tinyurl.com/p6zz9ex

1: socket = TSocket.TSocket('localhost', 9090)
2: transport = TTransport.TBufferedTransport(socket)
3: protocol = TBinaryProtocol.TBinaryProtocol(transport)
4: client = Calculator.Client(protocol)
5: transport.open()
6:
7: sum = client.add(1, 2)
8: print('1 + 2 = %d' % sum)
9:
10: transport.close()
```



```
# cat ~/repo/thrift_example/TARGETS

thrift_library(
    name = "myservice_thrift",
    languages = ["ruby", "php", "cpp", "python"],
    thrift_srcs = {"myservice.thrift": ["Calculator"]},
)

python_binary(
    name = "myservice_server",
    main_module = "thrift_example.server",
    srcs = ["server.py"],
    deps = [":myservice_thrift-py"],
)

python_binary(
    name = "myservice_client",
    main_module = "thrift_example.client",
    srcs = ["client.py"],
    deps = [":myservice_thrift-py"],
)
```

```
pallotron@dev:~/repo $ fbconfig thrift_example/
```

```
Configuring 20 targets
```

```
Writing 208/208 rules to makefile [100%]
```

```
Done
```

```
pallotron@dev:~/repo $ fbmake opt
```

```
  thrift/compiler/thriftl.ll
```

```
  common/memory/JEMalloc.cpp
```

```
  [...]
```

```
Linking _build/opt/thrift/compiler/thrift...
```

```
_build/opt/thrift/compiler/thrift
```

```
thrift_example/myservice.thrift
```

```
_build/opt/thrift_example/myservice_thrift-Calculator-pyremote.lpar
```

```
_build/opt/thrift_example/myservice_client.lpar
```

```
_build/opt/thrift_example/myservice_server.lpar
```

```
_build/opt/thrift_example/myservice_server.par
```

```
_build/opt/thrift_example/myservice_thrift-Calculator-pyremote.par
```

```
_build/opt/thrift_example/myservice_client.par
```



Python as a DSL

Enter configerator...

Thrift schema:



```
struct Job {  
  1: string name,  
  2: Scheduling  
    scheduling,  
}  
struct Scheduling {  
  1: int priority,  
}
```

Python config:

```
import Thrift(file)  
myConfig = Job(  
  name = "ajobname",  
  scheduling =  
    Scheduling(  
      priority = 3  
    )  
)  
export(myConfig)
```

Python validator:

```
def validate_job(job):  
  ensureNotEmpty(job.name)  
  ensurePositive(  
    job.scheduling.priority  
  )  
  
addValidator(myConfig,  
  validate_job)
```




```
1: def consume():
2:
3:     handle = dict()
4:     ctc = ConfiguratorThriftClient()
5:     fname = "example/config"
6:
7:     def on_update(name):
8:         handle['config'] = ctc.getJSONConfigContent(name)
9:
10:    ctc.startMonitoringConfig(fname)
11:    ctc.setUpdateCallback(on_update)
12:    on_update(fname)
13:
14:    return handle
```


Benefits of using configerator

Code review on configuration changes



<http://www.jasonawesome.com/wp-content/uploads/2010/06/sally-code-review-300x256.png>

Automated validation




http://sounddesignlive.com/wp-content/uploads/2013/08/sound-design-live-automate_all_the_things.jpeg

Clear representation of complex config



No service restart required to pick up changes



<http://social.microsoft.com/Forums/getfile/25914/>


Code review at FB

- Code **is** peer reviewed
- We use Phabricator
- Originally developed @ FB by Evan Priestley
- Go check www.phabricator.org
- Supports git/svn/hg



PHABRICATOR

- (Some) PEPs enforced and push blocking
- Common code dir for accelerated development
- Unit tests results are clearly visible in our code review tool



<http://memegenerator.net/Bill-Lumbergh-Office-Space>

Suggestions for reviewers:

- Reviewing code is not writing code
- Nitpicks are not enough
- Stop typing and have a conversation
- Build trust
- Find people to emulate

Making diffs easier to be reviewed:

- One and only one thesis
- Do No Harm
- Mix and match reviewers
- Write good test plans
- You are not your code

Production Python at Facebook

PROJECT KEANU

- Automated load tests of regions/clusters/hosts
- Implements feedback control loop
- No simulated traffic
- Goal is to find out “breaking point” == real capacity
- Runs daily
 - Peek time
 - When no incidents
 - Skips drained clusters



KEANU

Questions?

facebook